

Assignment 2: Implementing Feedforward neural networks with Keras and TensorFlow

Name: Zahir Mulla

```
In [12]: #Rolln0: B512012  
#Class: BE-IT(B)
```

```
In [1]: #installations  
from sklearn.preprocessing import LabelBinarizer  
from sklearn.metrics import classification_report  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense  
from tensorflow.keras.optimizers import SGD  
from tensorflow.keras.datasets import mnist  
from tensorflow.keras import backend as K  
import matplotlib.pyplot as plt  
import numpy as np
```

```
In [2]: #grabbing the mnist dataset  
((X_train, Y_train), (X_test, Y_test)) = mnist.load_data()  
X_train = X_train.reshape((X_train.shape[0], 28 * 28 * 1))  
X_test = X_test.reshape((X_test.shape[0], 28 * 28 * 1))  
X_train = X_train.astype("float32") / 255.0  
X_test = X_test.astype("float32") / 255.0
```

```
In [3]: lb = LabelBinarizer()  
Y_train = lb.fit_transform(Y_train)  
Y_test = lb.transform(Y_test)
```

```
In [4]: #building the model  
model = Sequential()  
model.add(Dense(128, input_shape=(784,), activation="sigmoid"))  
model.add(Dense(64, activation="sigmoid"))  
model.add(Dense(10, activation="softmax"))
```

```
In [5]: sgd = SGD(0.01)  
epochs=10  
model.compile(loss="categorical_crossentropy", optimizer=sgd, metrics=["accuracy"])  
H = model.fit(X_train, Y_train, validation_data=(X_test, Y_test), epochs=epochs, bat
```

```

Epoch 1/10
469/469 [=====] - 6s 8ms/step - loss: 2.2969 - accuracy:
0.1602 - val_loss: 2.2589 - val_accuracy: 0.2495
Epoch 2/10
469/469 [=====] - 3s 7ms/step - loss: 2.2349 - accuracy:
0.3125 - val_loss: 2.2058 - val_accuracy: 0.4334
Epoch 3/10
469/469 [=====] - 3s 7ms/step - loss: 2.1744 - accuracy:
0.4578 - val_loss: 2.1346 - val_accuracy: 0.5290
Epoch 4/10
469/469 [=====] - 3s 7ms/step - loss: 2.0897 - accuracy:
0.5428 - val_loss: 2.0319 - val_accuracy: 0.5811
Epoch 5/10
469/469 [=====] - 3s 7ms/step - loss: 1.9704 - accuracy:
0.5804 - val_loss: 1.8909 - val_accuracy: 0.6133
Epoch 6/10
469/469 [=====] - 3s 7ms/step - loss: 1.8142 - accuracy:
0.6122 - val_loss: 1.7176 - val_accuracy: 0.6407
Epoch 7/10
469/469 [=====] - 3s 7ms/step - loss: 1.6344 - accuracy:
0.6448 - val_loss: 1.5297 - val_accuracy: 0.6721
Epoch 8/10
469/469 [=====] - 4s 8ms/step - loss: 1.4520 - accuracy:
0.6755 - val_loss: 1.3516 - val_accuracy: 0.7004
Epoch 9/10
469/469 [=====] - 3s 7ms/step - loss: 1.2868 - accuracy:
0.7057 - val_loss: 1.1983 - val_accuracy: 0.7260
Epoch 10/10
469/469 [=====] - 4s 8ms/step - loss: 1.1487 - accuracy:
0.7325 - val_loss: 1.0740 - val_accuracy: 0.7516

```

```

In [6]: #making the predictions
predictions = model.predict(X_test, batch_size=128)
print(classification_report(Y_test.argmax(axis=1), predictions.argmax(axis=1), target

```

```

79/79 [=====] - 1s 5ms/step
              precision    recall  f1-score   support

     0         0.81         0.97         0.88         980
     1         0.78         0.99         0.87        1135
     2         0.82         0.69         0.75        1032
     3         0.62         0.90         0.73        1010
     4         0.67         0.80         0.73         982
     5         0.81         0.27         0.41         892
     6         0.84         0.87         0.85         958
     7         0.78         0.87         0.83        1028
     8         0.82         0.55         0.66         974
     9         0.67         0.54         0.60        1009

 accuracy         0.75        10000
 macro avg         0.76         0.74         0.73        10000
 weighted avg         0.76         0.75         0.74        10000

```

```

In [7]: #plotting the training loss and accuracy
plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, epochs), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, epochs), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, epochs), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, epochs), H.history["val_accuracy"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")

```

```
plt.ylabel("Loss/Accuracy")  
plt.legend()
```

Out[7]: <matplotlib.legend.Legend at 0x212d6c185e0>



In []: